

Secretary-Style Split Search for Decision Trees

Mathias Valla^a

^a*Chaire DIALog, Institut Louis Bachelier, Paris, France*

Abstract

Exhaustive CART split search can dominate training time on continuous tabular data because every admissible covariate–threshold pair is evaluated at each node. We study secretary-style split-search rules that inspect candidates online and stop according to split-level, nested, covariate-level, parametric, or blockwise rank-inspired criteria. Using 100 repeated runs on 113 regression datasets and 122 classification datasets from PMLB, we compare these rules with the exhaustive `scikit-learn` splitter and extremely randomized trees (ERT) baselines. After aligning exploration with ERT-style continuous threshold draws and optimizing exact selection, the simple split-level secretary and prophet rules preserve predictive accuracy much better than in earlier implementations but offer little runtime gain. The clearest practical compromises are instead covariate-level secretary search and ERT baselines, especially in classification. Secretary-style split search is therefore a family of task-dependent time–accuracy trade-offs rather than a universal replacement for exhaustive or ERT-style search.

Keywords: decision trees, CART, secretary problem, randomized split search, benchmark evaluation, computational efficiency

1. Introduction

Classification and regression trees remain widely used for tabular data because they combine flexibility, interpretability, and strong empirical performance [1]. Their standard induction rule is greedy: at each internal node, CART evaluates candidate covariate–threshold pairs and selects the split maximizing the impurity decrease. For continuous covariates, the number of admissible thresholds can be large, so exhaustive split search may dominate training time.

This raises a basic question: is it always necessary to evaluate all admissible splits in order to grow an accurate tree? In many nodes, only a small subset of candidates attains gains close to the optimum. Existing randomized tree procedures exploit this idea indirectly. Random forests subsample covariates [2], and extremely randomized trees (ERT) randomize threshold choice [3]. These methods, however, are not formulated as stopping rules on a stream of candidate gains.

We recast split search as a secretary-style online selection problem. In the classical secretary problem, items arrive sequentially in random order and the decision maker must stop online to select a highly ranked item [4, 5]. The classical skip rule, with asymptotic success probability $1/e$, shows that a short exploration phase can be followed by a simple acceptance rule with a nontrivial guarantee [4, 5]. More recent work has introduced rank-adaptive rules based on the availability of a single sample, including the block-rank family [6]. These ideas suggest a natural analogue for tree induction: if covariates and thresholds are explored in random order, split selection becomes a stopping problem on a stream of exact impurity gains.

We study four secretary-style split-search strategies for CART, ranging from split-level stopping to covariate-level stopping and model-based calibration, and also include a blockwise rank-inspired variant and a one-sample prophet-style baseline.

The contribution is empirical and methodological rather than theoretical. We benchmark the proposed rules against the standard exhaustive `scikit-learn` splitter and against ERT-style random-threshold baselines [3] on a large collection of regression and classification datasets from the PMLB benchmark suite [7]. The goal is to characterize the time–accuracy trade-off, quantify the variability induced by randomized split exploration, and identify which stopping rules offer the best practical compromise at the single-tree level before studying ensembles built from the same splitters.

Code availability. The implementation and the scripts used to reproduce the experiments are available at <https://github.com/MathiasValla/EarlyStoppingTrees>.

2. Methods

2.1. Split-search setting

Consider a node $\mathcal{N}$ containing n observations $\{(X_i, Y_i)\}_{i=1}^n$ with $X_i \in \mathbb{R}^p$. Let $\mathcal{J} = \{1, \dots, p\}$ be the set of covariates, and let $\mathcal{S}_j$ be the set of admissible

thresholds for covariate j . A candidate split is a pair (j, s) with $s \in \mathcal{S}_j$, producing children of sizes (n_L, n_R) . For an impurity I , the gain of split (j, s) is

$$G_{j,s} = I(\mathcal{N}) - \frac{n_L}{n}I(\mathcal{N}_L) - \frac{n_R}{n}I(\mathcal{N}_R). \quad (1)$$

For regression, I is the mean squared error impurity. For classification, we consider Gini and entropy impurities [1]. Once a candidate threshold is chosen, its gain is computed exactly on the full node data. In the optimized implementations studied here, randomness comes from the order in which covariates are visited and, for split-level exploration, from ERT-style continuous threshold draws within visited covariates.

To study early stopping, we impose a random order on visited covariates and use simple node-size-dependent exploration budgets. In the experiments, the exploration budget at a node is parameterized as a function of node size n , with schedules of order n/e , $0.1n$, $\log n$, and $\sqrt{n}$. The acceptance stage then combines these randomized exploration draws with exact within-covariate scans.

2.2. Stopping rules

Method S uses a random covariate order together with a split-level exploration budget. A subset of visited covariates is used for exploration; on each such covariate, a small number of ERT-style continuous thresholds is sampled and the best resulting gain updates the exploration maximum. Subsequent covariates are then scanned exactly, and the first covariate whose best exact split beats the exploration maximum is accepted. If no such covariate appears, the best exploration split is returned. This is the optimized tree-split analogue of the classical skip-then-accept secretary rule [4, 5].

Method S_{all} applies the secretary rule at the covariate level. For a visited covariate j , its score is

$$V_j = \max_{s \in \mathcal{S}_j} G_{j,s}.$$

Covariates are processed in random order, and a secretary rule is applied to the sequence (V_j) . When a covariate is accepted, the best threshold within that covariate is returned. This method is more expensive than S for each inspected covariate, but it preserves exact within-covariate optimization.

Method S^2 introduces a nested secretary structure. The outer rule is applied to covariates in random order. For each visited covariate j , an inner exploration phase samples a small number of ERT-style continuous thresholds,

and an exact within-covariate scan then returns the best threshold whose gain exceeds the inner exploration level. This yields a realized score $\tilde{V}_j$ for the covariate, and the outer rule is then applied to the sequence $(\tilde{V}_j)$. The outer stage selects the covariate and the inner stage selects the threshold.

Method S_{par} is a model-based variant. For a given covariate j , a small number of continuous random thresholds is first evaluated exactly, yielding gain samples used to fit a working distributional model or an empirical upper quantile. A covariate-specific acceptance threshold is then derived from this fit, and the covariate is represented by the first exact threshold whose gain exceeds the fitted level. The goal is to use a small number of exact evaluations to calibrate a conservative acceptance rule, especially in regression where the gain law is analytically tractable under Gaussian assumptions.

Inspired by the block-rank algorithm of Nuti and Vondrák [6], we also evaluate a blockwise rank-inspired heuristic. For a visited covariate, admissible thresholds are scanned in contiguous blocks of size about $\sqrt{n}$, each block is summarized by its exact best gain, and a secretary-style rule is applied to the resulting stream of block maxima. This is not intended as a literal transcription of the original block-rank schedule, but rather as a tree-specific blockwise analogue motivated by the same idea of using coarser rank information than full exhaustive search.

We also include a one-sample prophet-style baseline. For each visited covariate, one ERT-style continuous threshold is sampled and evaluated exactly, and the maximum of these gains is used as an empirical acceptance threshold. A second pass then scans the visited covariates exactly and returns the first covariate admitting a split above this threshold, with the original exploration split used as a fallback if no such covariate appears. This baseline is motivated by prophet inequalities from samples [8, 9].

2.3. Motivating heuristics

This subsection is motivational only. We do not attempt a formal analysis of tree induction under secretary-style exploration, and the paper’s contribution remains empirical and methodological.

The classical secretary problem shows that, under a random arrival order, a simple skip-then-accept rule with exploration length of order N/e is asymptotically optimal for selecting the best item [4, 5]. This is the only theoretical fact used directly in the benchmark design: it motivates the default n/e exploration schedule that we treat as the representative secretary-style setting in the main comparison.

The parametric variant S_{par} is motivated first by regression. Under homoscedastic Gaussian errors, the reduction in residual sum of squares associated with a fixed split is proportional to a χ_1^2 variable, so the corresponding regression gain admits a simple positive law after rescaling. This motivates fitting a working gain distribution from a small sample of exact evaluations and using it to calibrate a conservative acceptance threshold. For classification, no equally clean finite-sample law is available, but chi-square- or gamma-like approximations remain natural heuristics under standard asymptotics [10]. In the present paper, however, this idea is used only as a practical calibration device; we do not derive guarantees for the resulting stopping rule.

All methods were implemented in a common open-source package built around a modified tree splitter, with identical tree-growing hyperparameters across methods. This ensures that the observed differences are attributable to split-search rules rather than to other changes in the induction procedure.

3. Results

3.1. Experimental protocol

We compared secretary-style split search against the standard exhaustive splitter of `scikit-learn` and against ERT-style random-threshold baselines with `max_features` set to 1, $p/3$, $2p/3$, and p , on a broad collection of regression and classification datasets from the PMLB benchmark suite [7]. For each dataset, the same 5-fold cross-validation splits and tree hyperparameters were used across methods. Because the proposed procedures are randomized, each dataset–method pair was evaluated over 100 independent runs with different estimator seeds. After filtering to datasets that completed successfully under the common protocol, the final benchmark retained 113 regression datasets and 122 classification datasets for each impurity criterion.

The primary computational metric is wall-clock training time. To make the timing results more interpretable, we also track split-search effort through the number of gain evaluations, either measured directly for the secretary-style splitters or reconstructed deterministically a posteriori for the exhaustive and ERT baselines. Predictive performance is measured by RMSE in regression and by F1-score in classification; effort plots, reliability summaries, and additional visualizations are reported in the supplementary material. Results are normalized relative to the exhaustive splitter through a speedup factor and a relative loss metric. Because different splitters can induce different tree shapes, total tree-level effort can occasionally exceed the exhaustive

Table 1: Summary of the PMLB benchmark used in the study.

Task	N	median n	median p	$n_{\min}$ – $n_{\max}$	$p_{\min}$ – $p_{\max}$
Regression	113	500	10	47–40768	2–124
Classification (Gini)	122	736	13	32–58000	2–240
Classification (Entropy)	122	736	13	32–58000	2–240

baseline on a minority of datasets even when the node-level rule evaluates fewer gains on any given feature.

3.2. *Global trade-off*

Figure 1 summarizes the main Pareto comparison between speedup and predictive degradation. Each point corresponds to one dataset summarized by the median over the 100 runs, and the larger method markers summarize cross-dataset centroids. Size-stratified versions of the same Pareto comparison are reported in Supplementary Figure S1. Three conclusions are immediate: the time–accuracy trade-off depends strongly on the task; the strongest speed-oriented baselines remain the ERT family; and the optimized split-level secretary rules now lie much closer to the exhaustive baseline in predictive loss than in runtime.




Figure 1: Global Pareto comparison of exhaustive, secretary-style, and ERT-style split-search methods. Each point represents one dataset summarized by the median over 100 runs. The horizontal axis reports the median percentage of training time saved relative to the exhaustive `scikit-learn` splitter, and the vertical axis reports the median percentage loss relative to the same baseline. Left: regression with RMSE loss. Middle: classification with Gini impurity and F1 loss. Right: classification with entropy impurity and F1 loss. Large white-edged markers indicate cross-dataset centroids, and faint Pareto frontiers summarize the dataset-level and centroid-level envelopes.

Table 2 reports a representative benchmark summary. For the secretary families, we use the theoretically motivated n/e schedule; for the parametric strategy, we report the single representative $S_{\text{par}}(\sqrt{n}, q = 0.75)$, which remains Pareto-optimal in the screened summaries; and for ERT we report all four `max_features` budgets. In addition to centroid time saved and loss, the table reports gain-evaluation effort saved, which helps separate algorithmic budget reductions from residual implementation overhead.

The centroid time-saved column corresponds to the horizontal coordinate of the large markers in Figure 1. The differences between Figure 1 and Supplementary Figures S13–S14 should therefore be interpreted mainly as implementation effects: gain-evaluation effort can decrease substantially even when the realized wall-clock savings remain more modest.

The regression benchmark remains unfavorable to aggressive approximation, but the optimized split-level rules change the interpretation. The strongest speed-oriented baselines are still the ERT variants, which save about 54% to 71% of training time at substantial RMSE degradation. Among secretary-style rules, the clearest low-loss compromise is now $S_{\text{all}}(n/e)$, with

Table 2: Representative global summary by method after the final 100-run no-limit benchmark. The n/e schedule is used for the secretary families, $S_{\text{par}}(\sqrt{n}, q = 0.75)$ is the single parametric representative shown in the main comparison, and ERT denotes extremely randomized trees with four m_{try} budgets. Centroid time saved is the horizontal centroid of the dataset-level median points shown in Figure 1, expressed relative to the exhaustive `scikit-learn` splitter. Effort saved is the percentage reduction in gain evaluations relative to exhaustive split search. Loss is expressed in percent.

Task	Method	Centroid time saved (%)	Median effort saved (%)	Median loss (%)	90th pct. loss (%)	$P(\text{loss} \leq 5\%)$	$P(\text{speedup} \geq 20\%)$
Regression	Exhaustive	0.00	0.00	0.000	0.000	1.000	0.000
Regression	$S(n/e)$	15.67	46.85	3.881	10.546	0.593	0.442
Regression	$S^2(n/e)$	5.46	1.88	1.752	6.717	0.814	0.009
Regression	$S_{\text{all}}(n/e)$	16.44	17.42	1.812	6.642	0.832	0.425
Regression	Block-rank	60.46	90.76	31.907	55.316	0.159	0.947
Regression	1-sample prophet	30.13	58.38	6.432	15.475	0.434	0.743
Regression	ERT ($m_{\text{try}} = 1$)	71.05	98.93	36.675	60.558	0.115	0.965
Regression	ERT ($m_{\text{try}} = p/3$)	66.59	96.70	28.366	44.535	0.150	0.965
Regression	ERT ($m_{\text{try}} = 2p/3$)	60.50	92.95	15.232	28.983	0.221	0.956
Regression	ERT ($m_{\text{try}} = p$)	54.04	88.80	8.462	16.287	0.354	0.903
Classification (Gini)	Exhaustive	0.00	0.00	0.000	0.000	1.000	0.000
Classification (Gini)	$S(n/e)$	10.89	42.45	0.101	1.722	0.992	0.361
Classification (Gini)	$S^2(n/e)$	6.13	-1.10	0.045	1.091	0.992	0.025
Classification (Gini)	$S_{\text{all}}(n/e)$	12.74	15.27	0.033	1.072	0.992	0.328
Classification (Gini)	$S_{\text{par}}(\sqrt{n}, q=0.75)$	19.74	42.77	0.250	2.824	0.943	0.426
Classification (Gini)	Block-rank	45.89	71.89	2.675	14.638	0.639	0.738
Classification (Gini)	1-sample prophet	16.64	41.50	0.030	1.820	0.984	0.418
Classification (Gini)	ERT ($m_{\text{try}} = 1$)	54.19	96.34	4.573	21.931	0.525	0.803
Classification (Gini)	ERT ($m_{\text{try}} = p/3$)	49.04	89.25	1.736	7.923	0.779	0.779
Classification (Gini)	ERT ($m_{\text{try}} = 2p/3$)	43.69	84.10	1.006	3.803	0.934	0.738
Classification (Gini)	ERT ($m_{\text{try}} = p$)	38.14	84.57	0.291	2.598	0.967	0.623
Classification (Entropy)	Exhaustive	0.00	0.00	0.000	0.000	1.000	0.000
Classification (Entropy)	$S(n/e)$	12.59	41.95	0.100	1.490	1.000	0.410
Classification (Entropy)	$S^2(n/e)$	4.95	-1.93	0.069	1.214	1.000	0.008
Classification (Entropy)	$S_{\text{all}}(n/e)$	12.59	15.05	0.064	0.992	1.000	0.352
Classification (Entropy)	$S_{\text{par}}(\sqrt{n}, q=0.75)$	19.47	42.17	0.523	2.741	0.943	0.426
Classification (Entropy)	Block-rank	45.08	70.23	2.364	14.255	0.631	0.762
Classification (Entropy)	1-sample prophet	17.23	42.00	0.093	1.627	0.984	0.443
Classification (Entropy)	ERT ($m_{\text{try}} = 1$)	54.73	95.82	4.430	21.948	0.516	0.811
Classification (Entropy)	ERT ($m_{\text{try}} = p/3$)	49.27	88.55	1.802	7.676	0.779	0.787
Classification (Entropy)	ERT ($m_{\text{try}} = 2p/3$)	43.91	83.77	1.076	3.939	0.951	0.738
Classification (Entropy)	ERT ($m_{\text{try}} = p$)	38.61	83.46	0.486	2.303	0.975	0.631

centroid time saved 16.4% and median loss 1.81%. The optimized $S(n/e)$ rule and the one-sample prophet save 15.7% and 30.1% of time at median losses 3.88% and 6.43%, respectively. The earlier predictive gap of the split-level rules was therefore not purely algorithmic, but the new benchmark also shows that large reductions in gain evaluations do not automatically translate into proportional wall-clock savings.

The classification results remain more favorable than the regression results, but they also sharpen the distinction between quality-preserving and speed-preserving approximations. The optimized split-level secretary and prophet variants now achieve very small median losses, yet they still do not deliver the largest wall-clock gains. In both Gini and entropy classification, $S_{\text{all}}(n/e)$ is the strongest low-loss secretary-style compromise, with centroid time saved about 12.6% to 12.7% and median loss 0.03% to 0.06%. The optimized $S(n/e)$ rule saves 10.9% of time for Gini and 12.6% for entropy at about 0.10% median loss, while the one-sample prophet reaches 16.6% and 17.2% with median losses 0.03% and 0.09%. The selected parametric representa-

tive $S_{\text{par}}(\sqrt{n}, q = 0.75)$ remains Pareto-optimal in the screened summaries and saves about 19.5% to 19.7% of time, but it remains less reliable than $S_{\text{all}}(n/e)$. Supplementary Figure S2 still shows that S_{par} is most interesting on the smaller classification datasets rather than as a global default.

A plausible explanation for the strong task asymmetry observed in the benchmark is that classification impurities are more tolerant to approximate split search than regression impurity. For Gini and entropy, many nearby thresholds can induce very similar child class proportions and therefore nearly equivalent impurity reductions. In such cases, partial random exploration may still identify a split whose downstream predictive effect is close to that of the exhaustive optimum. By contrast, regression gains based on MSE appear to depend more sharply on the precise threshold location, because small changes in the partition can alter child means and variances continuously and hence affect the final predictions more directly. Supplementary Figure S12 examines this interpretation directly under exhaustive search by comparing the mass of thresholds within 5% and 10% of the nodewise optimum, the gap between the best gain and the median random-threshold gain, and the width of the near-optimal region on the winning covariate. The resulting comparison remains more nuanced than a single scalar notion of flatness: some diagnostics are larger for regression, whereas the best-versus-random gap is larger for classification. This suggests that the empirical advantage of classification is not explained by one universal notion of broader near-optimal threshold regions alone, and likely reflects a combination of impurity geometry, competition across covariates, downstream loss sensitivity, and the cost of exact replay.

3.3. Reliability and dataset structure

Figure 2 complements the centroid-based summaries by showing reliability under explicit operating constraints. The dominant pattern is not that faster methods are automatically reliable. Under the reliability criteria considered here and on the present benchmark, the most reliable secretary-style procedures are $S_{\text{all}}(n/e)$, then $S^2(n/e)$, and then $S(n/e)$, whereas block-rank and the ERT baselines achieve speed more aggressively but with much weaker reliability under strict loss constraints. In classification, $\text{ERT}(p)$ is consistently less reliable than those three secretary-style rules, and block-rank is substantially less reliable still. Size-stratified loss-only reliability curves are reported in Supplementary Figure S6, and the full grid of joint success probabilities over all datasets is reported in Supplementary Figure S7.




Figure 2: Reliability summaries for secretary-style split-search methods. Top row: empirical probability of achieving a predictive loss below a tolerance ε . Bottom panels: empirical probability of simultaneously achieving at least a prescribed time saving and a predictive loss below a prescribed tolerance. Separate panels are shown for regression, classification with Gini impurity, and classification with entropy impurity. Size-stratified and method-specific variants are reported in the supplementary material.

The reliability ranking therefore differs from the speed ranking. On the present benchmark, $S_{\text{all}}(n/e)$ is the most reliable method overall, followed by $S^2(n/e)$ and $S(n/e)$; these procedures keep losses under control much more consistently than block-rank or the ERT baselines. In regression, $S_{\text{all}}(n/e)$ and $S^2(n/e)$ form the strongest high-reliability tier, with $S(n/e)$ behind them. Block-rank remains interesting as a speed-oriented compromise, but not as a high-reliability default. Supplementary Figures S3–S5 report task-specific ridgeline summaries of within-dataset variability, and Supplementary Figures S9–S11 provide complementary whisker summaries of the same repeated-run variability.




Figure 3: Regime maps as a function of dataset size and dimensionality. Each point is a benchmark dataset represented in the $(\log_{10} n, \log_{10} p)$ plane and colored by the split-search strategy that achieves the largest median speedup among methods whose median loss does not exceed a fixed tolerance τ . Columns correspond to $\tau = 0.5\%, 1\%, 2.5\%$. Rows correspond to regression, classification with Gini impurity, and classification with entropy impurity.

Figure 3 shows that the usefulness of secretary-style split search depends on dataset structure. The colors do not identify the method with the smallest loss or the largest time saving separately; they identify the fastest method among those that satisfy the displayed loss tolerance. This is why exhaustive search can remain optimal in some regions despite zero loss, and why ERT with $m_{\text{try}} = 1$ is not always selected despite being the fastest raw baseline.

Classification tasks are clearly more favorable than regression tasks. Within classification, and under the tighter tolerances $\tau = 0.5\%$, 1% , and 2.5% , covariate-level secretary rules dominate larger portions of the regime map than split-level secretary rules. In regression, exhaustive split search remains the safest choice on much of the benchmark, and secretary-style rules are only competitive in more limited regions.

3.4. Sensitivity and negative result for the parametric strategy

We examined four exploration schedules based on node size: n/e , $0.1n$, $\log n$, and $\sqrt{n}$. The full results are reported in the supplementary material. Changing the schedule affects the quantitative trade-off more than the qualitative ranking. For S and S^2 , more aggressive schedules produce slightly larger wall-clock savings at the price of slightly worse predictive degradation. For S_{all} , the n/e schedule consistently provides the best compromise, especially in classification, which justifies its use as the default representative in the main text. The corresponding size-stratified Pareto analyses, reliability summaries, and effort comparisons are reported in Supplementary Figures S1, S6, S7, S13, and S14.

The parametric secretary method S_{par} was evaluated in an initial screening study over a grid of fitting sample sizes and quantile levels. Most configurations are dominated in practice: the cost of fitting the working model outweighs any benefit from early stopping, especially in regression and on larger classification datasets. The representative retained in the main text, $S_{\text{par}}(\sqrt{n}, q = 0.75)$, remains Pareto-optimal in the screened summaries and remains interesting primarily on smaller classification datasets. Even there, its global performance is less robust than the strongest ERT and S_{all} baselines. This mostly negative result shows that probabilistic calibration is not useful by itself when the fitting overhead is too large relative to the gain-computation cost.

For orientation, the supplementary material is organized as follows: Supplementary Figure S1 gives the size-stratified Pareto comparison; Supplementary Figure S2 gives the S_{par} screening study; Supplementary Figures S3–S5 give the ridgeline summaries; Supplementary Figures S6–S7 give the reliability and joint-success plots; Supplementary Figure S8 gives the tolerance-constrained quickest-method regime maps in the $(\log n, p/n)$ plane; Supplementary Figures S9–S11 give the within-dataset whisker summaries; Supplementary Figure S12 gives exhaustive gain-landscape diagnostics for the

regression–classification asymmetry; and Supplementary Figures S13–S14 compare split-search effort across methods and against wall-clock savings.

4. Conclusion, limitations, and outlook

This paper asked whether exhaustive CART split search can be replaced, at least in part, by secretary-style sequential selection without incurring an unacceptable loss in predictive performance. The answer depends strongly on both the stopping rule and the prediction task, and it also depends on implementation details. After aligning exploration with ERT-style continuous threshold draws and optimizing the exact-selection phase, the simple secretary and prophet variants preserve predictive performance much more closely than in the earlier implementation, but they do so with little wall-clock gain. The methods that still offer the clearest practical speedups are the ERT baselines and the covariate-level strategy S_{all} , whereas block-rank is better viewed as a speed-oriented but comparatively unreliable compromise.

The study reveals a clear task asymmetry. Regression trees are much less tolerant to incomplete split exploration than classification trees if one insists on real time savings. On the regression benchmarks considered here, the only low-loss secretary-style variants are effectively near-exhaustive in runtime, whereas the faster methods incur substantial loss. In classification, by contrast, the losses of several approximate methods remain modest enough to make the time–accuracy trade-off practically interesting. This suggests that the value of approximate split search depends both on the impurity criterion and on how efficiently a given stopping rule can convert reduced gain-evaluation effort into actual runtime savings.

A second important outcome is mostly negative but informative. The parametric secretary variant S_{par} , although statistically motivated by gain-distribution approximations, is not competitive globally once fitting overhead is taken into account. One low-overhead configuration remains interesting on small classification datasets, but overall the family is less robust than the strongest ERT and S_{all} baselines. This cautions against assuming that a principled probabilistic calibration will automatically lead to an effective split-search procedure.

The study is limited to single trees, benchmark datasets, and a deliberately small family of stopping rules, and is intended primarily as an empirical characterization rather than a complete theory. This restriction is deliberate: understanding the time–loss trade-off for a single tree is a necessary first step

before one can interpret what these splitters might contribute inside larger ensembles.

A natural next step is to study the same split-search rules in bagging and boosting, but this remains genuine future work rather than a claim of the present paper. Because split-search cost is incurred at every node and across many trees, even moderate node-level savings may accumulate into substantial end-to-end gains, and ensembling may damp the predictive effect of occasional suboptimal splits. At the same time, those ensemble-level effects cannot be interpreted cleanly without first understanding the single-tree regime studied here. A second direction is methodological: the performance of the block-rank rule suggests that adaptive rank-based or sample-based stopping rules may be preferable to the classical secretary rule in this context. Overall, secretary-style split search is not a universal replacement for exhaustive CART search, but in favorable settings, especially in classification, it offers a simple and interpretable way to trade split-search time against predictive performance.

References

- [1] L. Breiman, J. H. Friedman, R. A. Olshen, C. J. Stone, *Classification and Regression Trees*, Wadsworth, Belmont, CA, 1984.
- [2] L. Breiman, Random forests, *Machine Learning* 45 (1) (2001) 5–32. [doi:10.1023/A:1010933404324](https://doi.org/10.1023/A:1010933404324).
- [3] P. Geurts, D. Ernst, L. Wehenkel, Extremely randomized trees, *Machine Learning* 63 (1) (2006) 3–42. [doi:10.1007/s10994-006-6226-1](https://doi.org/10.1007/s10994-006-6226-1).
- [4] J. P. Gilbert, F. Mosteller, Recognizing the maximum of a sequence, *Journal of the American Statistical Association* 61 (313) (1966) 35–73. [doi:10.1080/01621459.1966.10502008](https://doi.org/10.1080/01621459.1966.10502008).
- [5] T. S. Ferguson, Who solved the secretary problem?, *Statistical Science* 4 (3) (1989) 282–296. [doi:10.1214/ss/1177012493](https://doi.org/10.1214/ss/1177012493).
- [6] P. Nuti, J. Vondrák, Secretary problems: The power of a single sample, in: *Proceedings of the 2023 ACM-SIAM Symposium on Discrete Algorithms*, SIAM, 2023, pp. 2015–2029. [doi:10.1137/1.9781611977554.ch77](https://doi.org/10.1137/1.9781611977554.ch77).

- [7] J. D. Romano, T. T. Le, W. L. Cava, J. T. Gregg, D. J. Goldberg, N. L. Ray, P. Chakraborty, D. Himmelstein, W. Fu, J. H. Moore, Pmlb v1.0: an open-source dataset collection for benchmarking machine learning methods, *Bioinformatics* 38 (3) (2022) 878–880. [doi:10.1093/bioinformatics/btab665](https://doi.org/10.1093/bioinformatics/btab665).
- [8] A. Rubinstein, J. Z. Wang, S. M. Weinberg, Optimal single-choice prophet inequalities from samples, in: 11th Innovations in Theoretical Computer Science Conference (ITCS 2020), Vol. 151 of LIPIcs, Schloss Dagstuhl – Leibniz-Zentrum für Informatik, 2020, pp. 60:1–60:10. [doi:10.4230/LIPIcs.ITCS.2020.60](https://doi.org/10.4230/LIPIcs.ITCS.2020.60).
- [9] E. Samuel-Cahn, Comparison of threshold stop rules and maximum for independent nonnegative random variables, *The Annals of Probability* 12 (4) (1984) 1213–1216. [doi:10.1214/aop/1176993150](https://doi.org/10.1214/aop/1176993150).
- [10] B. Laurent, P. Massart, Adaptive estimation of a quadratic functional by model selection, *The Annals of Statistics* 28 (5) (2000) 1302–1338. [doi:10.1214/aos/1015957395](https://doi.org/10.1214/aos/1015957395).